

Nullpointer's Glorious Intro to Buffer_Overflows

morn.

The Basics: What Is A Buffer?

- Holds input to program
- Fixed size
- Think arrays
 - `char buff[10]` is buffer of size 10
 - usually lack bounds checking (*not always*)

```
#include <stdio.h>

int main () {

    char buff[10];

    int overwriteMe = 0;

    scanf("%s",buff);

    if(overwriteMe != 0x1337){
        printf("Sorry, you failed\n");

        return 0;
    }
    else {
        printf("Success!\n");
    }

    return 0;
}
```

What goes on in the background

- Function calls create a frame on stack
 - Think of stack like “plates” of memory
 - holds buffer
 - holds vars
 - holds return addresses
- Buffer Overflow allows writing onto other parts of stack

So how does it work?

- Buffer has fixed size (usually)
- Other vars immediately follow buffer on stack
- Writing too much for unchecked buffer allows modifying the stack
- Easy to do with PYTHON :D

Important Stuff!

- Need to know how many bytes to deal with
- Addresses = 4 bytes, Ints = 4 bytes, Chars = 1 byte (language specific, look them up)
- 1 Byte = 2 Hex digits, e.g. 0xFF is 1 byte, 0xAABB is 2 bytes, etc.
- Different architectures use different “endianness” to read hex (next slide)
- Depending on compiler, stack smashing protection enabled by default
 - use compiler flag `-fno-stack-protector`

Very briefly on Endianness

- Order computers read hex
- two main kinds: consider 0xABCDEF12
 - Big Endian: human readable
 - 0xABCDEF12 -> 0xABCDEF12
 - Little Endian: Bytes are reversed
 - 0xABCDEF12 -> 0x12EFCDA B

The Code

- Take a look at example0 and example0.c
- Should be the following:

```
//Nullp0inter's Glorious intro to Buffer Overflows
//This one is just the basics, the buffer overfloweth with cyber riches yall

#include <stdio.h>

int main() {

    char buff[10];
    int overwriteMe = 0;

    scanf("%s",buff);

    if (overwriteMe == 0) {...
```

- You get the gist of it

So what does our stack look like?

The Stack:

- Call to function main builds a stack frame
 - pushes buff and other vars to stack
- int overwriteMe stored just below buffer
- *grows downward to lower addresses*
- Deals with hex representation
 - digits 0-9 and A-F (10-15)

Addresses	Contents
0xFFFFFA	buff[0]
0xFFFFF9	buff[1]
0xFFFFF8	buff[2]
⋮	⋮
0xFFFFF0	buff[9]
0xFFFFEF	overwriteMe = 0
0xFFFFEB	OtherStuff
0xFFFFE7	More Stuff
0xFFFFE3	Even More Stuff
0xFFFFDF	Probably a return address

So what happens when we input to the buffer?

- Buffer only holds so much data
 - If filled with more than it can hold, gets written past onto stack
- If we put 11 A's in a 10 char buffer, we write 1 byte past buffer
 - notice value of overwriteME

• They don't think it be like it do but it is

Addresses	Contents
0xFFFFFA	buff[0] = A (0x41)
0xFFFFF9	buff[1] = A (0x41)
0xFFFFF8	buff[2] = A (0x41)
⋮	⋮
0xFFFFF0	buff[9] = A (0x41)
0xFFFFEF	overwriteMe = A (0x41)
0xFFFFEB	OtherStuff
0xFFFFE7	More Stuff
0xFFFFE3	Even More Stuff
0xFFFFDF	Probably a return address

So what happens when we input to the buffer?

-Continued-

- overwriteMe overwritten with 11th A
- Ascii / Unicode gets converted to hex to be stored
- Poor buffer means near full control of stack with overflows

Addresses	Contents
0xFFFFFA	buff[0] = A (0x41)
0xFFFFF9	buff[1] = A (0x41)
0xFFFFF8	buff[2] = A (0x41)
⋮	⋮
0xFFFFF0	buff[9] = A (0x41)
0xFFFFEF	overwriteMe = A (0x41)
0xFFFFEB	OtherStuff
0xFFFFE7	More Stuff
0xFFFFE3	Even More Stuff
0xFFFFDF	Probably a return address

- Malicious user can exploit control of stack to do as they please
 - bypass checks
 - escape program
 - get shell
- All of this is great but how exactly can it be done? There are a couple of methods but my favorite is....

PYTHON

yayyyyyy lmao !!!

Using Python for Input

- can be done in single command and piped to program
 - Piping of the form: `<command> | <second command> | <third command> ... etc`
- Can reverse “ and ‘ as long as used consistently

The specifics

- In general it starts like this:

```
python -c "print 'A' * <buffer size>  
+ <any extra stuff needed>" |  
./<program binary>
```

- “Why don’t I just type everything manually or generate it and copy and paste?” Two main reasons:
 - buffer not always nice and small
 - not all hex has a character to it

example0

- requires overwriting integer overwriteMe
- goal is to make overwriteMe *anything* other than 0
- not so bad

example1

- In example1.c, notice overwriteMe checked against 0x1337
- Need to write past 10 char buffer with a *specific value* (0x1337)
- Take a minute to try it on your own

Solving example1

- Open a terminal (if you have not done so already) and type the command:

```
python -c "print 'A' * 10 + '\x37\x13'"  
| ./example1
```

- So what is going on there?
 - buffer filled with 10 A's (buff = AAAAAAAAAAAA)
 - Hex value concatenated to A's and passed in Little Endian

Hints for example2

- example2 requires overwriting a function pointer
- figure out on your own
- use hints below to figure it out
 - `gdb`
 - `b main`
 - `p <function name>`

Stack frame precedence

- Example in testing0 and testing0.c
- Typically buffer always first in stack
 - followed by vars in *reverse* order they were declared
 - set to print values, mess around a bit

Extremely Useful Tools

- Several extremely useful tools for overflow, ROP, and reversing:
 - GDB - GNU debugger to step through execution
 - peda - python based gdb extension to show assembly, registers, stack, etc during execution (nice colors)
 - radare2 - disassembler and more, get it from GitHub and *not* from apt repos (repo build severely out of date)
 - - can generate patterns as well (ragg2 -P <len> -r)

Resources

- [gera's insecure programming](#)
- [Hacking the Art of Exploitation](#)
- [Smashing the Stack for fun and profit](#) by aleph1 (Phrack Mag)